

# Assignment No. 1

**Predict the price of the Uber ride from a given pickup point to the agreed drop-off location. Perform following tasks:**

- 1. Pre-process the dataset.**
- 2. Identify outliers.**
- 3. Check the correlation.**
- 4. Implement linear regression and random forest regression models.**
- 5. Evaluate the models and compare their respective scores like R2, RMSE, etc**

## Step 1: Download the dataset from Kaggle

- 1. Dataset link:** [Uber Fares Dataset](#)
  2. Go to the Kaggle dataset page.
  3. Click on the **Download** button to download the uber.csv file.
  4. After downloading, make sure the dataset is placed in the same directory where you plan to open and run your Jupyter Notebook.
    - If the dataset is located in a different directory, update the file path in the code where `pd.read_csv()` is used (e.g., `pd.read_csv('path/to/uber.csv')`).
- 

## Step 2: Open Jupyter Notebook

- 1. Open Jupyter Notebook:**
    - Launch Jupyter Notebook from your system (either through **Anaconda** or by typing `jupyter notebook` in the command line).
    - This will open the Jupyter environment in your web browser.
  - 2. Create a new notebook:**
    - Once Jupyter opens, navigate to the directory where you placed the dataset.
    - Click **New** -> **Python 3** to create a new Python notebook.
- 

## Step 3: Paste the code into Jupyter Notebook

Now you will start adding the code in chunks. Here's how each step of the code works.

### 3.1: Import necessary libraries

```
import pandas as pd
```

```
import numpy as np

import matplotlib.pyplot as plt

import seaborn as sns

from sklearn.model_selection import train_test_split
```

- **Explanation:** This step imports essential libraries:
    - pandas for handling data and creating DataFrames.
    - numpy for numerical operations.
    - matplotlib and seaborn for visualization.
    - train\_test\_split from sklearn.model\_selection to split the data into training and testing sets.
  - **Execute this code block:** Place this code in the first cell of your notebook and run it by pressing **Shift + Enter**.
- 

### 3.2: Load the dataset

```
df = pd.read_csv('uber.csv')

df.head()
```

- **Explanation:**
    - pd.read\_csv('uber.csv') reads the dataset into a pandas DataFrame.
    - df.head() displays the first 5 rows of the dataset so you can verify it loaded correctly.
  - **Note:** If the dataset is saved in another folder, you need to provide the full path, like pd.read\_csv('C:/path/to/uber.csv').
- 

### 3.3: Check dataset shape

```
df.shape
```

- **Explanation:** This line checks the number of rows and columns in the dataset.
    - The result will be something like (n\_rows, n\_columns).
- 

### 3.4: Check for missing values

```
df.isnull()
```

- **Explanation:** This line checks for missing values (null values) in the dataset. It will return True for each cell that contains a missing value.

---

### 3.5: Drop unnecessary columns

```
df.drop(columns=["Unnamed: 0", "key"], inplace=True)

df.head()
```

- **Explanation:**
  - `df.drop()` removes columns that are not useful for analysis.
  - `inplace=True` ensures that the changes are made directly to the DataFrame without needing to reassign it.
  - After dropping the columns, we use `df.head()` again to confirm that the columns are removed.

---

### 3.6: Count missing values in each column

```
df.isnull().sum()
```

- **Explanation:** This counts the total number of missing values in each column. This is important for determining which columns need further cleaning.

---

### 3.7: Fill missing values in specific columns

```
df['dropoff_latitude'].fillna(value=df['dropoff_latitude'].mean(), inplace=True)

df['dropoff_longitude'].fillna(value=df['dropoff_longitude'].median(), inplace=True)
```

- **Explanation:**
  - This fills missing values in the `dropoff_latitude` column using the column's mean value.
  - The `dropoff_longitude` column is filled using the column's median value. These choices (mean and median) are typical strategies for handling missing data.
  - `inplace=True` ensures that the changes are made directly to the DataFrame.

---

### 3.8: Check the data types of each column

```
df.dtypes
```

- **Explanation:** This displays the data types (e.g., integer, float, object) of each column. This is useful to identify if any column needs conversion, especially datetime columns.
-

### 3.9: Convert 'pickup\_datetime' to datetime format

```
df.pickup_datetime = pd.to_datetime(df.pickup_datetime)

df.dtypes
```

- **Explanation:** This converts the pickup\_datetime column to a proper datetime format, allowing for easier manipulation and feature extraction.
    - After conversion, you can check the data types again to confirm the change.
- 

### 3.10: Create new columns based on 'pickup\_datetime'

```
df = df.assign(hour=df.pickup_datetime.dt.hour,
               day=df.pickup_datetime.dt.day,
               month=df.pickup_datetime.dt.month,
               year=df.pickup_datetime.dt.year,
               dayofweek=df.pickup_datetime.dt.dayofweek)
```

- **Explanation:**
    - This extracts various components (hour, day, month, year, day of the week) from the pickup\_datetime column and adds them as new columns to the DataFrame.
- 

### 3.11: Drop the 'pickup\_datetime' column

```
df = df.drop(["pickup_datetime"], axis=1)

df
```

- **Explanation:** Since we've extracted useful information from pickup\_datetime, we no longer need the original column, so it's dropped.
- 

### 3.12: Define a function to calculate travel distance

```
from math import *

def distance_formula(longitude1, latitude1, longitude2, latitude2):

    travel_dist = []

    for pos in range(len(longitude1)):

        lon1, lan1, lon2, lan2 = map(radians, [longitude1[pos], latitude1[pos],
        longitude2[pos], latitude2[pos]])
```

```

dist_lon = lon2 - lon1

dist_lan = lan2 - lan1

a = sin(dist_lan/2)**2 + cos(lan1) * cos(lan2) * sin(dist_lon/2)**2

c = 2 * asin(sqrt(a)) * 6371

travel_dist.append(c)

return travel_dist

```

- **Explanation:**

- This function calculates the distance between two geographical points using the **Haversine formula**, which accounts for the curvature of the Earth.
  - It converts the latitude and longitude values from degrees to radians and then applies trigonometry to compute the distance in kilometers.
- 

### 3.13: Add the calculated travel distance to the DataFrame

```

df['dist_travel_km'] = distance_formula(df.pickup_longitude.to_numpy(),
df.pickup_latitude.to_numpy(), df.dropoff_longitude.to_numpy(),
df.dropoff_latitude.to_numpy())

```

- **Explanation:** This uses the `distance_formula` function to calculate the distance between the pickup and dropoff points for each row in the DataFrame, and stores the result in the new `dist_travel_km` column.
- 

### 3.14: Define features (X) and target (y)

```

df_x =
df[['pickup_longitude','pickup_latitude','dropoff_longitude','dropoff_latitude','passenger_count','hour','day','month','year','dayofweek','dist_travel_km']]

df_y = df['fare_amount']

```

- **Explanation:**

- `df_x` includes the input features (longitude, latitude, passenger count, etc.) used for predicting the target.
  - `df_y` is the target variable (fare amount) that we want to predict.
- 

### 3.15: Split the data into training and testing sets

```

x_train, x_test, y_train, y_test = train_test_split(df_x, df_y, test_size=0.2,
random_state=1)

```

- **Explanation:**
    - `train_test_split` splits the data into training (80%) and testing (20%) sets.
    - `random_state=1` ensures that the split is reproducible.
- 

### 3.16: Train and predict using Linear Regression

```
from sklearn.linear_model import LinearRegression  
reg = LinearRegression()  
reg.fit(x_train, y_train)  
y_pred_lin = reg.predict(x_test)  
print(y_pred_lin)
```

- **Explanation:**
    - This initializes a linear regression model and trains it using the training data.
    - After training, it predicts the fare amounts for the test set and prints the predictions.
- 

### 3.17: Train and predict using Random Forest

```
from sklearn.ensemble import RandomForestRegressor  
rf = RandomForestRegressor(n_estimators=100)  
rf.fit(x_train, y_train)  
y_pred_rf = rf.predict(x_test)  
print(y_pred_rf)
```