

Assignment No. 4

Implement K-Nearest Neighbors algorithm on diabetes.csv dataset. Compute confusion matrix, accuracy, error rate, precision and recall on the given dataset.

Step 1: Download the dataset from Kaggle

1. **Dataset link:** [Diabetes Dataset](#)
 2. Go to the Kaggle dataset page.
 3. Download the diabetes.csv file.
 4. Save the dataset in the directory where you will run the Jupyter notebook.
-

Step 2: Open Jupyter Notebook

1. **Open Jupyter Notebook:**
 - Launch **Jupyter Notebook**.
 - Navigate to the directory where you saved the diabetes.csv file.
 - Create a new Python notebook.
-

Step 3: Import necessary libraries

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
```

- **Explanation:** These libraries will help load the dataset, split the data, and apply the KNN classifier.
-

Step 4: Load and inspect the dataset

```
# Load the dataset
diabetes_data = pd.read_csv('diabetes.csv')
# Display the first 5 rows of the dataframe
diabetes_data.head()
```

- **Explanation:** The dataset is loaded into a pandas DataFrame `diabetes_data`, and `head()` displays the first 5 rows for inspection.

```
diabetes_data.shape
```

- **Explanation:** This outputs the shape of the dataset, i.e., the number of rows and columns.

```
diabetes_data.describe()
```

- **Explanation:** This provides summary statistics of the numerical columns in the dataset.
-

Step 5: Define features (X) and target (Y)

```
# Drop the 'Outcome' column from the feature set
```

```
X = diabetes_data.drop(columns='Outcome', axis=1)
```

```
# Display the first 5 rows of X
```

```
X.head()
```

- **Explanation:** X contains all the features except the 'Outcome' column, which represents the target.

```
# Define the target variable
```

```
Y = diabetes_data['Outcome']
```

- **Explanation:** Y contains the target variable, 'Outcome', which we are trying to predict (whether the patient has diabetes or not).
-

Step 6: Split the data into training and testing sets

```
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=1)
```

- **Explanation:** The dataset is split into training and testing sets using `train_test_split()`, with 80% of the data used for training and 20% for testing. `random_state=1` ensures the split is consistent every time the code is run.
-

Step 7: Train and apply the KNN classifier

```
# Initialize and train a KNN classifier with 7 neighbors
```

```
KN = KNeighborsClassifier
```

```
knn = KN(n_neighbors=7)
```

```
knn.fit(x_train, y_train)
```

```
# Make predictions on the test set
```

```
y_pred = knn.predict(x_test)
```

```
print("Prediction: \n")
```

```
print(y_pred)
```

- **Explanation:**

- The KNN classifier is initialized with 7 neighbors (`n_neighbors=7`), meaning it will consider the 7 nearest neighbors to classify each test data point.
 - The model is trained on `x_train` and `y_train`, and predictions are made on `x_test`.
-

Step 8: Evaluate the model performance

1. Confusion Matrix

```
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score,  
recall_score
```

```
# Compute the confusion matrix
```

```
conf_matrix = confusion_matrix(y_test, y_pred)
```

```
print("Confusion Matrix:\n", conf_matrix)
```

- **Explanation:**

- The confusion matrix is computed using `confusion_matrix()`. It shows how well the model predicted the classes:
 - **True Negatives (TN):** Number of cases correctly classified as not having diabetes.
 - **False Positives (FP):** Number of cases incorrectly classified as having diabetes.
 - **False Negatives (FN):** Number of cases incorrectly classified as not having diabetes.
 - **True Positives (TP):** Number of cases correctly classified as having diabetes.

2. Accuracy

```
# Compute accuracy
```

```
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

3. Error Rate

```
# Compute error rate
error_rate = 1 - accuracy
print("Error Rate:", error_rate)
```

4. Precision

```
# Compute precision
precision = precision_score(y_test, y_pred)
print("Precision:", precision)
```

5. Recall

```
# Compute recall
recall = recall_score(y_test, y_pred)
print("Recall:", recall)
```